

CS463/516

Lecture 16


Segmentation as a graph problem

Segmentation as a graph problem

- Pixels and voxels, and relationship between them representable as graph
- Adjacency represented by edges connecting nodes
- Each node carries local features (intensity, gradient, or some texture measure)
- Can map 2d and 3d images on graph where scene elements are nodes and neighborhood is expressed by edges connecting the nodes
- Assigning weights to edges that represent local properties of a good segmentation allows finding a segmentation using graph optimization techniques
- Two such techniques are:
 - 1) minimum cost graph cuts (will cover this)
 - 2) minimum cost paths (will not cover this)
- Segmentation then either
 - a) produces closed path that encloses a segment, or
 - b) produces labeling of nodes so that all elements belonging to segment have same label


Graph cuts

- A *graph cut* creates two or more disconnected subgraphs by removing edges from a connected graph
- Graph cuts well suited to solve foreground segmentation problem
- *Minimum cost graph cut* consists of a set of edges such that the sum of the edge weights in the cut is minimal for all possible cuts
- Finding minimum cut is solved as a maximum flow problem.
 - Each edge assigned a flow capacity
 - All nodes $p \in P$ representing pixels are connected to two special nodes: source S and sink T via *terminal links*
 - Edges connecting neighboring pixels are called *neighbor links* (n -links), the set of all n -links is N
 - Assume that water is flowing from source, through graph, to the sink



Graph cuts

- Assume we know the segmentation, then weights for all links can be preset to describe the segmentation
- T-links from source to foreground pixels and from sink to background pixels are labeled '0', remaining t-link are labeled '1'
- n -links connecting foreground or background pixels to each other attributed a '1', and n -links connecting foreground pixel with background pixel receive '0' as attribute
- Optimal cut consists of all t-links and n-links with value of 0, since total cost for all edges in this case would be zero
- Set of edges of the cut consists of:
 - all edges connecting foreground pixels with the source,
 - all edges connecting background pixels with the sink
 - All edges at region boundaries between foreground and background pixels
- Hence, edges to one of the terminal nodes in the cut define the segment label of the segmentation
- Edges of the cut that are n-links represent boundary between segments



Graph cuts for image

segmentation: all pixels are connected to their neighbors via n -links and to a source and sink via t -links. The cut separates the source from the sink. All t -links that are part of the cut indicate association of a pixel either to foreground (t -link to source) or to the background.




Graph cuts

- Computing segmentation that is already known not particularly useful
- Strategy from previous slide can be used if assignment of pixel to foreground or background is less clear
- Instead of attributing pixels with certainty of belonging to foreground, $R_p(\text{frg}) = 0$, $R_p(\text{bkg}) = 1$, or background $R_p(\text{frg}) = 1$ and $R_p(\text{bkg}) = 0$, probabilities of a pixel of belonging to foreground or background are used to assign weights to t -links.
- If probabilities P for some pixel \mathbf{p} with coordinates (p_1, p_2, \dots) are $P(\mathbf{p} \in \text{frg})$ and $P(\mathbf{p} \in \text{bkg})$ then a possible measure suggested by inventors of foreground graph cuts for segmentation is:
 - $R_p(\text{frg}) = -\ln P(\mathbf{p} \in \text{frg})$ and $R_p(\text{bkg}) = -\ln P(\mathbf{p} \in \text{bkg})$
- Probabilities P specified based on domain knowledge (such as bone in CT ranges from 60 to 3000 HU), or can be estimated from histogram analysis of training data



Graph cuts

- $R_p(\text{frg}) = -\ln P(\mathbf{p} \in \text{frg})$ and $R_p(\text{bkg}) = -\ln P(\mathbf{p} \in \text{bkg})$

- Value of this measure is close to zero if probability of a pixel belonging to foreground or background is high and increases with decreasing probability



Example: flow values from source to pixels (foreground) and from pixels to sink (background) are inversely proportional to the probability of these pixels belonging to foreground or background




Graph cuts

- If nothing else about segment membership known, all n -links would receive constant weight, and application of graph-cut algorithm would be similar to histogram-based thresholding
- Power of graph-cut stems from combining local region attribute with boundary attributes
- Generally, segment boundaries indicated by change of intensity or texture
- Given that function $f(\mathbf{p})$ exists for each pixel \mathbf{p} that combines these three attributes in a proper manner for segmentation, difference $\|f(\mathbf{p}) - f(\mathbf{q})\|$ for two pixels \mathbf{p} and \mathbf{q} connected by n -link may serve as weight for this link

Large value for σ means small intensity differences counted as noise and do not receive a low flow value

- Can use following measure:
- $$B_{p,q} = \exp\left(-\frac{\|f(\mathbf{p}) - f(\mathbf{q})\|^2}{2\sigma^2}\right) \cdot \frac{1}{\|\mathbf{p} - \mathbf{q}\|}$$
 - Which resembles a non-normalized gaussian. σ is set so it separates high-frequency and high-amplitude noise from true edge variation, which is assumed to have lower amplitude in the high frequency range than noise.




Graph cuts


	Edge type	Weights	Edges
a)	n -link	$B_{p,q} = \exp\left[-\frac{\ f(\mathbf{p})-f(\mathbf{q})\ ^2}{2\sigma^2}\right] \cdot \frac{1}{\ \mathbf{p}-\mathbf{q}\ }$	all $(p, q) \in N$
	t -link	$R(\mathbf{p}, S) = \lambda \cdot (-\ln P(\mathbf{p} \in \text{"frg"}))$	all $(S, p), p \in P$
	t -link	$R(\mathbf{p}, T) = \lambda \cdot (-\ln P(\mathbf{p} \in \text{"bkg"}))$	all $(T, p), p \in P$

- If n -link weights used with all t -links having same weights, method produces results similar to watershed transform (which finds all boundaries at zero crossings)
- Combining the two measures compensates for deficiencies of one of the measures by another measure, weights for all links summarized in (a)
- Finding cut C separating S and T that minimizes cost of all edges in C minimizes:
- $E(C) = \lambda \cdot (\sum_{p \in S_C} R(\mathbf{p}, S) + \sum_{p \in T_C} R(\mathbf{p}, T)) + \sum_{(p,q) \in N_C} B_{p,q}$
 - Where N_C is subset of all edges in N that are part of cut C , and S_C and T_C are subsets of edges (\mathbf{p}, S) and (\mathbf{p}, T) , respectively, that are part of C .
- λ governs relative influence of boundary term with respect to the region term
 - Small λ used if probability distributions of region attribute overlap or are very flat (b)
- Segmentation could still succeed if sufficiently high number of n -links on segment boundary have low values (c)

b) Interactive input for foreground and background produces two distributions that overlap to such a large extent that a segmentation based on intensity alone will not produce desired grouping into a single foreground object surrounded by background



c) Intermediate range of values for λ results in the desired segmentation




Graph cuts

- Graph cuts have been extended further to include prior knowledge about region membership
- If foreground or background membership of some pixels known, their t -link weights can be set accordingly, to ensure minimum cost cut does not cut the corresponding link
- t -link weights for known foreground pixels \mathbf{p}_{frg} and background pixels \mathbf{p}_{bkg} are:
- $R(\mathbf{p}_{frg}, S) = 0 \wedge R(\mathbf{p}_{frg}, T) = 1 + \max_{p \in P} (\sum_{q: (p,q) \in N} B_{p,q})$, and
- $R(\mathbf{p}_{bkg}, T) = 0 \wedge R(\mathbf{p}_{bkg}, S) = 1 + \max_{p \in P} (\sum_{q: (p,q) \in N} B_{p,q})$
- Will always be costlier to cut t -link from T to a \mathbf{p}_{frg} or from S to a \mathbf{p}_{bkg} than to cut any n -link leading to pixel \mathbf{p}_{frg} and \mathbf{p}_{bkg} , respectively

Ford-Fulkerson algorithm

- Computation of min cut done by computing maximum flow between source and sink
- Flow from source to sink limited by a set of edges which are saturated with flow
- This set of edges forms closed boundary separating source from sink
- Iterative Ford-Fulkerson algorithm computes maximum flow
 - Augments flow until saturation (a)
- Algorithm keeps copy of $G_{residual}$ of graph $G = \{V, E\}$ with nodes $V = \{v_i\}$ and edges $E = \{e_j\}$.
- Weights $w_{res}(e_i)$ of edges in $G_{residual}$ are residual flow capacity, which at initialization is equal to capacity $w(e_i)$ of edges in G


a) Ford-Fulkerson algorithm find sequence of edges with maximum flow capacity at each step. Total flow is increased accordingly and remaining capacity is reduced



Ford-Fulkerson algorithm

- At each iteration, minimum cost path from S to T determined in $G_{residual}$ (a)
- Flow from S to T augmented by capacity of edge with smallest weight w_{min} along this path
- Residual capacity along edges of path is reduced by w_{min}
- Edges E with $w_{res}(e) = 0$ are removed and are part of the graph cut
- Process continues until T becomes unreachable from S
- Set of edges removed from residual capacity graph constitutes minimum cost graph cut


Algorithm is efficient in worst-case analysis for arbitrary graphs, but average performance for computing graph cuts for segmentation can be improved using Boykov's algorithm



Adding domain knowledge

- Basic graph cut segmentation produces segments based on homogeneity in the segments and boundary length
- Performance can be enhanced if additional domain knowledge is integrated into cost function
- 'blob' component has been used to separate heart from background in CT
- Blob component penalizes direction deviations of a line segment pq of a cut to a line from a prespecified blob center c to the location of p (a,b)
- this is a way to promote convex structures around a pre-specified blob center (which in this case is a location in center of heart specified by user)


Figure (a,b) – blob constraint penalizes cutting through edges between pixels the more the orientation varies from the orientation of a line between the blob center and the pixel. Cutting the edge depicted in (a) is cheaper than cutting the edge depicted in (b)



Normalized graph cuts

- Graph cuts can produce undesired results because the minimum-cut-maximum-flow tends to cut as few edges as possible (a)
- Can be helpful to add size constraint into equation that requires segments to have similar size
- *Normalized graph cut* segments an image into regions with one of two labels ***a*** or ***b*** with minimal costs
- Costs consist of original graph cut costs $cut(\mathbf{a}, \mathbf{b})$ weighted by association costs $assoc(\mathbf{a}, \mathbf{v})$ and $assoc(\mathbf{b}, \mathbf{v})$ between nodes of a segment to all nodes \mathbf{v} in the scene
- Total cost of a normalized cut is then $NCut(\mathbf{a}, \mathbf{b}) = \frac{cut(\mathbf{a}, \mathbf{b})}{assoc(\mathbf{a}, \mathbf{v})} + \frac{cut(\mathbf{a}, \mathbf{b})}{assoc(\mathbf{b}, \mathbf{v})}$

a) Graph cut segmentation may stop too early because inhomogeneous foreground on the right results in a weak data constraint. Hence, smoothness overtakes, which causes a short boundary close to the initialization



Normalized graph cuts

- Computing optimal cut is NP-complete, so we approximate:
- Problem mapped on a linear equation system representing graph as a matrix and unknown labeling by an indicator vector
- Indicator vector \mathbf{x} for nodes $\mathbf{N} = \{\mathbf{n}_1, \dots, \mathbf{n}_M\}$ in the graph has elements x_i as follows:
 - $x_i = \begin{cases} 1, & \text{if } \mathbf{n}_i \in A \\ -1, & \text{if } \mathbf{n}_i \in B \end{cases}$
- In $M \times M$ weighted matrix, \mathbf{W} denotes costs of the edge between nodes n_i and n_j
- The value of an element w_{ij} is $w_{ij} = \begin{cases} e(i, j), & \text{if } \mathbf{n}_i \text{ and } \mathbf{n}_j \text{ are connected by edge} \\ 0, & \text{otherwise} \end{cases}$
- Furthermore, diagonal matrix \mathbf{D} is created with entries d_{ii} representing association costs of a node \mathbf{n}_i to all other nodes: $d_{ii} = \sum_{j=1, M} w_{ij}$
- Cost function can then be written as:
- $NCut(A, B) = NCut(\mathbf{x}) = \frac{\sum_{\{(i,j)\}, x_i > 0, x_j < 0} -w_{ij}x_ix_j}{\sum_{\{i\}, x_i > 0} d_{ii}} + \frac{\sum_{\{(i,j)\}, x_i < 0, x_j > 0} -w_{ij}x_ix_j}{\sum_{\{i\}, x_i < 0} d_{ii}}$

Normalized graph cuts continued

- Introducing cost ratio k for the association costs of scene elements \mathbf{a} to the total association cost:

$$k = \frac{\sum_{\{i\}, x_i > 0} d_{ii}}{\sum_i d_{ii}}$$

- And setting $\mathbf{y} = (\mathbf{1} + \mathbf{x}) - \frac{k}{1-k}(\mathbf{1} - \mathbf{x})$
- It can be shown that cost of a cut specified by indicator vector \mathbf{x} is given by:
- $NCut(\mathbf{x}) = \frac{\mathbf{y}^T (\mathbf{D} - \mathbf{W}) \mathbf{y}}{\mathbf{y}^T \mathbf{D} \mathbf{y}}$
- Finding an optimal cut requires to find vector \mathbf{y} that minimizes $NCut(\mathbf{x})$
- Solutions can be generated by solving associated eigenproblem for:
$$(\mathbf{D} - \mathbf{W})\mathbf{y} = \lambda \mathbf{D} \mathbf{y}$$
- Solution for largest eigenvalue (i.e. the associated eigenvector) produces the vector \mathbf{y} and consequently the vector \mathbf{x} to create the labeling
- Result only approximates optimal solution, since \mathbf{x} may contain values that are neither -1 or 1
- In final step, real-valued vector is mapped onto an indicator vector
- Advantage of NCut is that it is an unsupervised method
 - This may be disadvantage, as object-specific information (e.g. location of organ) not part of segmentation model

Normalized graph cut example

Compute original seg

Do graph cut

Get gray matter segments
(hand-picked)

```
from skimage import data, segmentation, color
from skimage.future import graph
from matplotlib import pyplot as plt
from sklearn.cluster import MiniBatchKMeans

labels1 = segmentation.felzenszwalb(img,min_size=50)

g = graph.rag_mean_color(img, labels1, mode='similarity')
labels2 = graph.cut_normalized(labels1, g,max_edge=1)
plt.imshow(labels2)




labs = [108,22,45,14]
graymatter = np.zeros(img.shape)
for lab in labs:
    graymatter += labels2==lab
```

raw

original segmentation

normalized graph cuts

gray matter segmentation




Graph-based segmentation conclusion

- Mapping image on graph is straightforward, as are the definitions of global and local criteria for a good segmentation by node costs
- Methods presented show different ways to compute optimal segmentation based on cost functions using standard algorithms on graphs
- Advantage of graphs is their independence from dimensionality, and versatility to define different cost functions that represent different segmentation goals
- Much more to graph-based segmentation that we didn't cover:
 - Graph cuts to approximate Bayesian segmentation
 - Problem-specific adaptation of t -link weights
 - Segmentation as a *path problem*
 - Fuzzy connectedness
 - Image foresting transform
 - Random walks

Frangi filter

- Frangi filter used on raw image (a) to detect vessel or tube-like structures (b)

a)



b)



Hessian matrix: square matrix of second order partial derivatives of scalar field.

- Describes local curvature of function of many variables
- For 3D input image I , Hessian matrix is a 3×3 matrix composed of second order partial derivatives of I :

$$\nabla^2 I = \begin{bmatrix} \frac{\partial^2 I}{\partial x^2} & \frac{\partial^2 I}{\partial x \partial y} & \frac{\partial^2 I}{\partial x \partial z} \\ \frac{\partial^2 I}{\partial y \partial x} & \frac{\partial^2 I}{\partial y^2} & \frac{\partial^2 I}{\partial y \partial z} \\ \frac{\partial^2 I}{\partial z \partial x} & \frac{\partial^2 I}{\partial z \partial y} & \frac{\partial^2 I}{\partial z^2} \end{bmatrix}$$